

# Chapter 8 Code

*Xi Liang*

*5/29/2017*

## Contents

Example 1. Identifying frequently purchased groceries with association rules . . . . .	1
Data Preparation - creating a sparse matrix for transaction data . . . . .	1
Visualizing item support - item frequency plots . . . . .	2
Visualizing the transaction data - plotting the sparse matrix . . . . .	3
Training a model on the data . . . . .	4
Evaluating model performance . . . . .	6
Improving model performance . . . . .	6
Sorting the set of association rules . . . . .	6
Taking subsets of association rules . . . . .	7
Saving association rules to a file or data frame . . . . .	7

## Example 1. Identifying frequently purchased groceries with association rules

Our market basket analysis will utilize the purchase data collected from one month of operation at a real-world grocery store. The data contains 9,835 transactions or about 327 transactions per day.

### Data Preparation - creating a sparse matrix for transaction data

```
library(arules)

## Loading required package: Matrix
##
## Attaching package: 'arules'
## The following objects are masked from 'package:base':
##
##   abbreviate, write

groceries <- read.transactions("data/groceries.csv", sep = ",")
summary(groceries)

## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 169 columns (items) and a density of 0.02609146
##
## most frequent items:
##   whole milk other vegetables      rolls/buns      soda
##      2513      1903      1809      1715
##   yogurt      (Other)
##    1372      34055
##
## element (itemset/transaction) length distribution:
## sizes
##   1    2    3    4    5    6    7    8    9   10   11   12   13   14   15
```

```
## 2159 1643 1299 1005 855 645 545 438 350 246 182 117 78 77 55
## 16 17 18 19 20 21 22 23 24 26 27 28 29 32
## 46 29 14 14 9 11 4 6 1 1 1 1 3 1
##
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 1.000 2.000 3.000 4.409 6.000 32.000
##
## includes extended item information - examples:
## labels
## 1 abrasive cleaner
## 2 artif. sweetener
## 3 baby cosmetics
```

```
inspect(groceries[1:5])
```

```
## items
## [1] {citrus fruit,
## margarine,
## ready soups,
## semi-finished bread}
## [2] {coffee,
## tropical fruit,
## yogurt}
## [3] {whole milk}
## [4] {cream cheese,
## meat spreads,
## pip fruit,
## yogurt}
## [5] {condensed milk,
## long life bakery product,
## other vegetables,
## whole milk}
```


```
itemFrequency(groceries[, 1:5])
```

```
## abrasive cleaner artif. sweetener baby cosmetics baby food
## 0.0035587189 0.0032536858 0.0006100661 0.0001016777
## bags
## 0.0004067107
```

## Visualizing item support - item frequency plots


Eight items in the data with at least 10% support.

```
itemFrequencyPlot(groceries, support = 0.1)
```



Top 20 items in the data


```
itemFrequencyPlot(groceries, topN = 20)
```



Visualizing the transaction data - plotting the sparse matrix

To visualize the entire sparse matrix using `image()`

```
image(groceries[1:5])
```




From the diagram above, we observe that there are 5 rows and 169 columns, indicating 5 transactions and 169 possible items we requested.

We can also see that first, fourth, and fifth transactions contained four items each, and row three, five, two and four have an item in common.

Visualizing random selection of 100 transactions

```
image(sample(groceries, 100))
```



## Training a model on the data

We will attempt to use the default settings of support = 0.1 and confidence = 0.8.

```
apriori(data = groceries)
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.8    0.1    1 none FALSE              TRUE     5     0.1    1
## maxlen target  ext
```

```
##      10 rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE FALSE TRUE 2 TRUE
##
## Absolute minimum support count: 983
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [8 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 done [0.00s].
## writing ... [0 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

## set of 0 rules
```

We ended with 0 rules returned when we used the default settings, which is not surprising. Because support = 0.1 by default, in order to generate a rule, an item must have appeared in at least  $0.1 * 9,385 = 938.5$  transactions. Since only eight items appeared this frequently in our data, it's no wonder that we did not find any rules.

One way to approach the above problem of setting a min support threshold is to think about the smallest number of transactions you would need before you would consider a pattern interesting. For example, you could argue that if an item is purchased twice a day (about 60 times in a month of data), it may be an interesting pattern. From there, it is possible to calculate the support level needed to find only the rules matching at least that many transactions. Since  $60/9835 = 0.006$ , we'll try setting the support there first.

We will start with confidence threshold of 0.25, which means that in order to be included in the results, the rule has to be correct at least 25 percent of the time. We'll also set minlen = 2 to eliminate rules that contain fewer than two items.

```
groceryrules <- apriori(groceries, parameter = list(support = 0.006, confidence = 0.25, minlen = 2))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.25    0.1    1 none FALSE          TRUE      5  0.006      2
## maxlen target   ext
##      10 rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE FALSE TRUE 2 TRUE
##
## Absolute minimum support count: 59
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [109 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [463 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
groceryrules
```

```
## set of 463 rules
```

Our `groceryrules` object contains a set of 463 association rules. To determine whether any of them are useful, we will have to dig deeper.

## Evaluating model performance

```
summary(groceryrules)
```

```
## set of 463 rules
##
## rule length distribution (lhs + rhs):sizes
##   2   3   4
## 150 297  16
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   2.000  2.000   3.000   2.711   3.000   4.000
##
## summary of quality measures:
##      support      confidence      lift
##   Min.   :0.006101   Min.   :0.2500   Min.   :0.9932
##   1st Qu.:0.007117   1st Qu.:0.2971   1st Qu.:1.6229
##   Median :0.008744   Median :0.3554   Median :1.9332
##   Mean   :0.011539   Mean   :0.3786   Mean   :2.0351
##   3rd Qu.:0.012303   3rd Qu.:0.4495   3rd Qu.:2.3565
##   Max.   :0.074835   Max.   :0.6600   Max.   :3.9565
##
## mining info:
##      data ntransactions support confidence
##   groceries          9835    0.006      0.25
```

In our rule set, 150 rules have only two items, while 297 have three, and 16 have four.

We will inspect the first 3 rules in the `groceryrules` object:

```
inspect(groceryrules[1:3])
```

```
##      lhs      rhs      support      confidence lift
## [1] {pot plants} => {whole milk} 0.006914082 0.4000000 1.565460
## [2] {pasta}      => {whole milk} 0.006100661 0.4054054 1.586614
## [3] {herbs}      => {root vegetables} 0.007015760 0.4312500 3.956477
```

## Improving model performance

### Sorting the set of association rules

Depending upon the objects of the market basket analysis, the most useful rules might be the ones with the highest support, confidence, or lift. The best five rules according to the lift statistic can be examined using the following command:

```
inspect(sort(groceryrules, by = "lift")[1:5])
```

```
##      lhs      rhs      support      confidence lift
## [1] {herbs}      => {root vegetables} 0.007015760 0.4312500 3.956477
```

```
## [2] {berries}          => {whipped/sour cream} 0.009049314 0.2721713 3.796886
## [3] {other vegetables,
##      tropical fruit,
##      whole milk}    => {root vegetables} 0.007015760 0.4107143 3.768074
## [4] {beef,
##      other vegetables} => {root vegetables} 0.007930859 0.4020619 3.688692
## [5] {other vegetables,
##      tropical fruit}  => {pip fruit} 0.009456024 0.2634561 3.482649
```

### Taking subsets of association rules

Suppose that given the preceding rule, the marketing team is excited about the possibilities of creating an advertisement to promote berries, which are now in season. Before finalizing the campaign, however, they ask you to investigate whether berries are often purchased with other items. To answer this question, we will need to find all the rules that include berries in some form.

```
berryrules <- subset(groceryrules, items %in% "berries")
inspect(berryrules)
```

```
##      lhs      rhs      support    confidence lift
## [1] {berries} => {whipped/sour cream} 0.009049314 0.2721713 3.796886
## [2] {berries} => {yogurt} 0.010574479 0.3180428 2.279848
## [3] {berries} => {other vegetables} 0.010269446 0.3088685 1.596280
## [4] {berries} => {whole milk} 0.011794611 0.3547401 1.388328
```

### Saving association rules to a file or data frame

```
write(groceryrules, file = "groceryrules.csv",
      sep = ",", quote = TRUE, row.names = FALSE)
```

```
groceryrules_df <- as(groceryrules, "data.frame")
```

```
str(groceryrules_df)
```

```
## 'data.frame':   463 obs. of  4 variables:
## $ rules      : Factor w/ 463 levels "{baking powder} => {other vegetables}",...: 340 302 207 206 208 ...
## $ support    : num  0.00691 0.0061 0.00702 0.00773 0.00773 ...
## $ confidence: num  0.4 0.405 0.431 0.475 0.475 ...
## $ lift       : num  1.57 1.59 3.96 2.45 1.86 ...
```