

Chapter 9 Code

Xi Liang

5/29/2017

Contents

Example - Finding teen market segments using k-means clustering	1
Exploring and preparing the data	1
Data preparation - dummy coding missing values	2
Data preparation - imputing the missing values	3
Training a model on the data	3
Evaluating model performance	4
Improving model performance	5

Example - Finding teen market segments using k-means clustering

Exploring and preparing the data

```
teens <- read.csv("data/snsdata.csv")
str(teens)
```

```
## 'data.frame':   30000 obs. of  40 variables:
## $ gradyear      : int  2006 2006 2006 2006 2006 2006 2006 2006 2006 2006 ...
## $ gender        : Factor w/ 2 levels "F","M": 2 1 2 1 NA 1 1 2 1 1 ...
## $ age           : num  19 18.8 18.3 18.9 19 ...
## $ friends       : int  7 0 69 0 10 142 72 17 52 39 ...
## $ basketball    : int  0 0 0 0 0 0 0 0 0 0 ...
## $ football      : int  0 1 1 0 0 0 0 0 0 0 ...
## $ soccer        : int  0 0 0 0 0 0 0 0 0 0 ...
## $ softball      : int  0 0 0 0 0 0 0 1 0 0 ...
## $ volleyball    : int  0 0 0 0 0 0 0 0 0 0 ...
## $ swimming      : int  0 0 0 0 0 0 0 0 0 0 ...
## $ cheerleading  : int  0 0 0 0 0 0 0 0 0 0 ...
## $ baseball      : int  0 0 0 0 0 0 0 0 0 0 ...
## $ tennis        : int  0 0 0 0 0 0 0 0 0 0 ...
## $ sports        : int  0 0 0 0 0 0 0 0 0 0 ...
## $ cute          : int  0 1 0 1 0 0 0 0 0 1 ...
## $ sex           : int  0 0 0 0 1 1 0 2 0 0 ...
## $ sexy          : int  0 0 0 0 0 0 0 1 0 0 ...
## $ hot           : int  0 0 0 0 0 0 0 0 0 1 ...
## $ kissed        : int  0 0 0 0 5 0 0 0 0 0 ...
## $ dance         : int  1 0 0 0 1 0 0 0 0 0 ...
## $ band          : int  0 0 2 0 1 0 1 0 0 0 ...
## $ marching      : int  0 0 0 0 0 1 1 0 0 0 ...
## $ music         : int  0 2 1 0 3 2 0 1 0 1 ...
## $ rock          : int  0 2 0 1 0 0 0 1 0 1 ...
## $ god           : int  0 1 0 0 1 0 0 0 0 6 ...
## $ church        : int  0 0 0 0 0 0 0 0 0 0 ...
## $ jesus         : int  0 0 0 0 0 0 0 0 0 2 ...
## $ bible         : int  0 0 0 0 0 0 0 0 0 0 ...
```

```
## $ hair      : int  0 6 0 0 1 0 0 0 0 1 ...
## $ dress     : int  0 4 0 0 0 1 0 0 0 0 ...
## $ blonde    : int  0 0 0 0 0 0 0 0 0 0 ...
## $ mall      : int  0 1 0 0 0 0 2 0 0 0 ...
## $ shopping  : int  0 0 0 0 2 1 0 0 0 1 ...
## $ clothes   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ hollister : int  0 0 0 0 0 0 2 0 0 0 ...
## $ abercrombie : int  0 0 0 0 0 0 0 0 0 0 ...
## $ die       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ death     : int  0 0 1 0 0 0 0 0 0 0 ...
## $ drunk     : int  0 0 0 0 1 1 0 0 0 0 ...
## $ drugs     : int  0 0 0 0 1 0 0 0 0 0 ...
```

The data include 30,000 teenagers with four variables indicating personal characteristics and 36 words indicating interests.

From the brief overview of the data, feature **gender** contains some missing information, we want to know how many missing data there are in this feature, since knowing the sexuality of the individuals is important to this study.

```
prop.table(table(teens$gender, useNA = "ifany"))
```

```
##
##          F          M        <NA>
## 0.7351333 0.1740667 0.0908000
```

Around 9% have missing gender data, and interestingly, there are over four times as many females as males in the SNS data, suggesting that males are not as inclined to use SNS website as females.

Beside **gender**, we also find there are a lot of missing data in **age**.

```
summary(teens$age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##   3.086  16.312  17.287  17.994  18.259 106.927    5086
```

A total of 5,086 records (17%) have missing ages. Also concerning is the fact that the minimum and maximum values seem to be unreasonable. To ensure that these extreme values don't cause problems for the analysis, we will need to clean them up before moving on.

A more reasonable range of ages for the high school students includes those who are at least 13 years old and not yet 20 years old.

```
teens$age <- ifelse(teens$age >= 13 & teens$age <20,
                  teens$age, NA)
```

```
summary(teens$age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##   13.03  16.30  17.27  17.25  18.22  20.00    5523
```

Data preparation - dummy coding missing values

Since there is a large portion of the data contain missing values, instead of removing those values, we will create another factor beside male and female.

```
teens$female <- ifelse(teens$gender == "F" &
                      !is.na(teens$gender), 1, 0)

teens$no_gender <- ifelse(is.na(teens$gender), 1, 0)
```

```
table(teens$gender, useNA = "ifany")
```

```
##  
##      F      M <NA>  
## 22054  5222  2724
```

```
table(teens$female, useNA = "ifany")
```

```
##  
##      0      1  
##  7946 22054
```

```
table(teens$no_gender, useNA = "ifany")
```

```
##  
##      0      1  
## 27276  2724
```

Data preparation - imputing the missing values

```
mean(teens$age, na.rm = TRUE)
```

```
## [1] 17.25243
```

```
aggregate(data = teens, age ~ gradyear, mean, na.rm = TRUE)
```

```
##   gradyear    age  
## 1    2006 18.65586  
## 2    2007 17.70617  
## 3    2008 16.76770  
## 4    2009 15.81957
```

```
ave_age <- ave(teens$age, teens$gradyear, FUN = function(x) mean(x, na.rm = TRUE))
```

```
teens$age <- ifelse(is.na(teens$age), ave_age, teens$age)
```

```
summary(teens$age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
##   13.03   16.28   17.24   17.24   18.21   20.00
```

Training a model on the data

The `kmeans()` function requires a data frame containing only numeric data and a parameter specifying the desired number of clusters. We will be including all the interests in the data frame

```
interests <- teens[5:40]
```

```
interests_z <- as.data.frame(lapply(interests, scale))
```

We will try using $k = 5$ and see where it leads us.

```
set.seed(2345)
```

```
teens_clusters <- kmeans(interests_z, 5)
```

Evaluating model performance

```
teens_clusters$size
```

```
## [1] 871 600 5981 1034 21514
```

Here, we see the five clusters we requested. The smallest cluster has 600 teenagers while the largest cluster has 21,514. Although the large gap between the number of people in the largest and smallest cluster is slightly concerning, without examining these groups more carefully, we will not know whether or not this indicates a problem.

For a more in-depth look at the clusters, we can examine the coordinates of the cluster centroids using the `teens_clusters$centers` component

```
teens_clusters$centers
```

```
##      basketball      football      soccer      softball      volleyball      swimming
## 1  0.16001227  0.2364174  0.10385512  0.07232021  0.18897158  0.23970234
## 2 -0.09195886  0.0652625 -0.09932124 -0.01739428 -0.06219308  0.03339844
## 3  0.52755083  0.4873480  0.29778605  0.37178877  0.37986175  0.29628671
## 4  0.34081039  0.3593965  0.12722250  0.16384661  0.11032200  0.26943332
## 5 -0.16695523 -0.1641499 -0.09033520 -0.11367669 -0.11682181 -0.10595448
##      cheerleading      baseball      tennis      sports      cute
## 1  0.3931445  0.02993479  0.13532387  0.10257837  0.37884271
## 2 -0.1101103 -0.11487510  0.04062204 -0.09899231 -0.03265037
## 3  0.3303485  0.35231971  0.14057808  0.32967130  0.54442929
## 4  0.1856664  0.27527088  0.10980958  0.79711920  0.47866008
## 5 -0.1136077 -0.10918483 -0.05097057 -0.13135334 -0.18878627
##      sex      sexy      hot      kissed      dance      band
## 1  0.020042068  0.11740551  0.41389104  0.06787768  0.22780899 -0.10257102
## 2 -0.042486141 -0.04329091 -0.03812345 -0.04554933  0.04573186  4.06726666
## 3  0.002913623  0.24040196  0.38551819 -0.03356121  0.45662534 -0.02120728
## 4  2.028471066  0.51266080  0.31708549  2.97973077  0.45535061  0.38053621
## 5 -0.097928345 -0.09501817 -0.13810894 -0.13535855 -0.15932739 -0.12167214
##      marching      music      rock      god      church      jesus
## 1 -0.10942590  0.1378306  0.05905951  0.03651755 -0.00709374  0.01458533
## 2  5.25757242  0.4981238  0.15963917  0.09283620  0.06414651  0.04801941
## 3 -0.10880541  0.2844999  0.21436936  0.35014919  0.53739806  0.27843424
## 4 -0.02014608  1.1367885  1.21013948  0.41679142  0.16627797  0.12988313
## 5 -0.11098063 -0.1532006 -0.12460034 -0.12144246 -0.15889274 -0.08557822
##      bible      hair      dress      blonde      mall      shopping
## 1 -0.03692278  0.43807926  0.14905267  0.06137340  0.60368108  0.79806891
## 2  0.05863810 -0.04484083  0.07201611 -0.01146396 -0.08724304 -0.03865318
## 3  0.22990963  0.23612853  0.39407628  0.03471458  0.48318495  0.66327838
## 4  0.08478769  2.55623737  0.53852195  0.36134138  0.62256686  0.27101815
## 5 -0.06813159 -0.20498730 -0.14348036 -0.02918252 -0.18625656 -0.22865236
##      clothes      hollister      abercrombie      die      death
## 1  0.5651537331  4.1521844  3.96493810  0.043475966  0.09857501
## 2 -0.0003526292 -0.1678300 -0.14129577  0.009447317  0.05135888
## 3  0.3759725120 -0.0553846 -0.07417839  0.037989066  0.11972190
## 4  1.2306917174  0.1610784  0.26324494  1.712181870  0.93631312
## 5 -0.1865419798 -0.1557662 -0.14861104 -0.094875180 -0.08370729
##      drunk      drugs
## 1  0.035614771  0.03443294
## 2 -0.086773220 -0.06878491
```

```
## 3 -0.009688746 -0.05973769
## 4  1.897388200  2.73326605
## 5 -0.087520105 -0.11423381
```

Improving model performance

```
teens$cluster <- teens_clusters$cluster
```

After assigning cluster numbers back to the data, we would like to see how the cluster assignment relates to individual characteristics. For example, here is the personal information for the first five teens in the SNS data:

```
teens[1:5, c("cluster", "gender", "age", "friends")]
```

```
##   cluster gender    age friends
## 1      5      M 18.982        7
## 2      3      F 18.801         0
## 3      5      M 18.335        69
## 4      5      F 18.875         0
## 5      4  <NA> 18.995        10
```

We can also look at the demographic characteristics of the clusters.

```
aggregate(data = teens, age ~ cluster, mean)
```

```
##   cluster    age
## 1      1 16.86497
## 2      2 17.39037
## 3      3 17.07656
## 4      4 17.11957
## 5      5 17.29849
```

We can see that mean age does not vary much by cluster.

On the other hand, there are some substantial differences in the proportion of females by gender.

```
aggregate(data = teens, female ~ cluster, mean)
```

```
##   cluster  female
## 1      1 0.8381171
## 2      2 0.7250000
## 3      3 0.8378198
## 4      4 0.8027079
## 5      5 0.6994515
```

This is a very interesting finding as we did not use gender data to create the clusters. Cluster 1 and 3 are nearly 84% female, and these clusters show above the mean interest level on all fashion/shopping related topics. While Cluster 2 and Cluster 5 are only 70% female.

We suspect that the clusters are predictive of the number of friends the users have:

```
aggregate(data = teens, friends ~ cluster, mean)
```

```
##   cluster friends
## 1      1 41.43054
## 2      2 32.57333
## 3      3 37.16185
## 4      4 30.50290
```

5 5 27.70052